<u>A review of some evaluation heuristics for the game of Isolation</u> By Ezenwere I.K Emmanuel

<u>Custom score</u>: the custom score uses the probability of winning as an evaluation heuristic for board states. When compared to AB_Improved it performed better with an average of 70.29% win rate while AB_Improved had a 68.29% Win rate. It was a slightly increased performance. Besides having the best performance compared to the rest, I choose this because I'm convinced that having a stochastic heuristic function is the right way to go, especially because we can't be 100% certain of the outcome of a game without searching to end-game.

Also I found it beautiful the consistency this function gives when we eventually search to end game, here's how:

Assuming we search to end game, we have S moves and our opponent is isolated and has no more moves left, P(winning) = S/(S+0) = 1

While if we are isolated and our opponent has S moves then P(Winning) = 0/(0+S) = 0

And that just makes logical sense.

Here's the program:

```
def custom_score(game, player):
    if game.is_loser(player):
        return float("-inf")

if game.is_winner(player):
    return float("inf")

own_moves = len(game.get_legal_moves(player))
    opp_moves = len(game.get_legal_moves(game.get_opponent(player)))
    win_probability = float(own_moves / (own_moves + opp_moves))

return win_probability
```

<u>Custom score 2</u>: the custom score 2 square(my_moves) - square(opp_moves), ,this performed better than AB_Improved on average with an average win rate of 70.65% while AB_Improved had 68.29%. The motivation behind this heuristic was to magnify number of moves of the respective players in other to have a greater contrast between relative board advantage of each player.

```
if game.is_loser(player):
    return float("-inf")

if game.is_winner(player):
    return float("inf")

own_moves = len(game.get_legal_moves(player))
    opp_moves = len(game.get_legal_moves(game.get_opponent(player)))

return float(own_moves**2 - opp_moves**2)
```

<u>Custom score 3</u>: the custom score 3 probability of winning / (my_moves - opp_moves), ,this performed slightly less than AB_Improved on average with an average win rate of 62.14% while AB_Improved had 68.29%. The motivation behind this was to come up with a heuristic that uses some form of probability and the same algorithm for AB_Improved, (my_moves - opp_moves). This didn't perform so much as expected.

```
def custom_score_3(game, player):
    if game.is_loser(player):
        return float("-inf")

if game.is_winner(player):
    return float("inf")

own_moves = len(game.get_legal_moves(player))
opp_moves = len(game.get_legal_moves(game.get_opponent(player)))
score1 = float(own_moves / (own_moves + opp_moves))
score2 = float(own_moves - opp_moves)

if score1 == 0:
    return float("inf")
else:
    return score2 / score1
```

Win rate 1	AB_Improved	AB_Custom	AB_Custom_2	AB_Custom_3
1	<u>58.6</u>	<u>61.4</u>	<u>65.7</u>	<u>60.0</u>
2	<u>68.6</u>	<u>64.3</u>	70.0	<u>65.7</u>
<u>3</u>	<u>74.3</u>	72.9	<u>78.6</u>	<u>54.3</u>
4	<u>55.7</u>	70.0	<u>77.1</u>	<u>67.1</u>
<u>5</u>	<u>67.1</u>	<u>77.1</u>	<u>65.7</u>	<u>55.7</u>
<u>6</u>	<u>71.4</u>	<u>74.3</u>	<u>71.4</u>	<u>61.4</u>
7	<u>67.1</u>	<u>68.6</u>	<u>67.1</u>	<u>58.6</u>
8	<u>71.4</u>	<u>68.6</u>	<u>67.1</u>	<u>62.9</u>
9	<u>75.7</u>	<u>71.4</u>	72.9	<u>71.4</u>
<u>10</u>	72.9	74.3	70.0	64.3
Average win rate after 10	68.29	70.29	70.56	62.14

<u>batches</u>

Remark:

I'm convinced that the right way to go about this is using deep learning; here's the proposed architecture, by using a deep neural network to generate some initial random function that maps an NxN (board dimension) matrix representation of the board to a single float, next using this same random function as our heuristic function we play the game a 1000 times, with cost function $C = (100 - win_rate)$ reflecting the percentage of wins, next we back propagate and repeat this process until we can get C to converge to 0.

At this point we will have a representation that most accurately maps any given board state to the most accurate evaluation score.

I spent a majority of my time on this project designing such a neural network but I couldn't implement this eventually because I couldn't find out how to express backpropagation in this scenario where I don't have labels and input features beforehand.

I'm still working on this and I'm currently studying the Alpha Go paper to understand if and how they handled this challenge.